



# AlbaTablut


*A “rising” tablut player ...still rising*

# The approach

**Goal:** attempt to replicate the idea behind AlphaGo and AlphaGo Zero, in order to learn more about how it works!

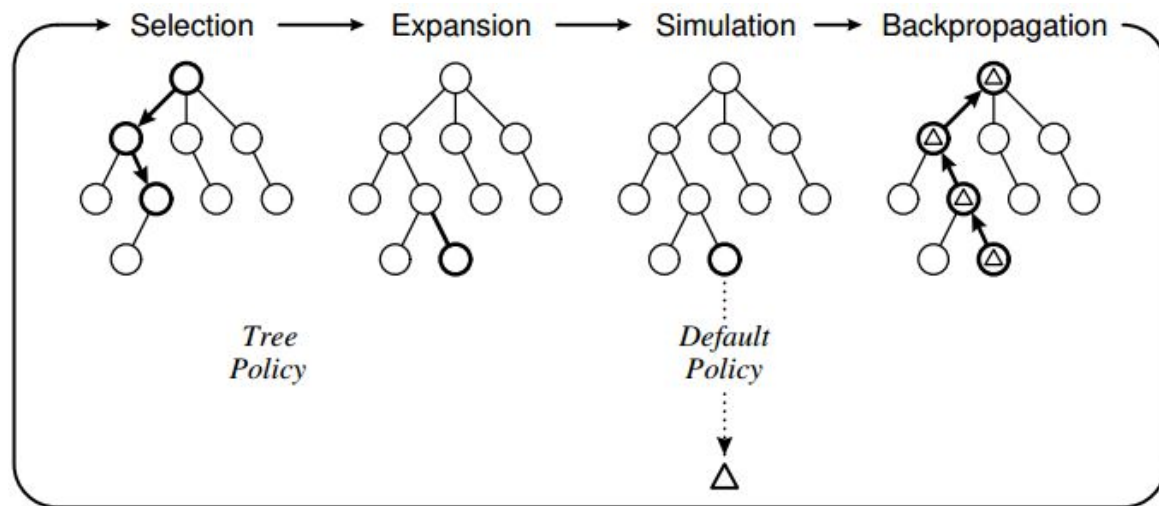
**Method:**

- Monte Carlo Tree Search to explore the search space
- Neural networks to improve MCTS, trained using past years' games data



# Search space exploration - Monte Carlo Tree Search

- Explore in greater depth the paths deriving from **more promising actions**
- **Evaluate** how good a newly expanded node is by simulating the game evolution
- Backpropagate the knowledge up the tree for the next exploration round




# Improvements from Neural Networks

## Policy network

**Problem:** Expansion and Selection (initially) are random, uninformed → inefficient exploration


Returns a **distribution over actions**, highlighting the more promising ones, given a board state



## Value network

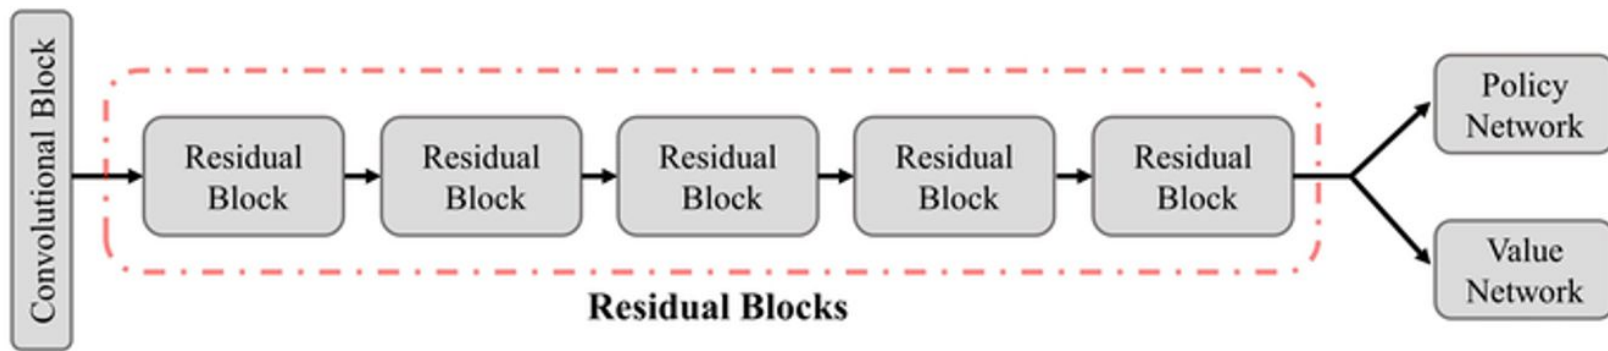
**Problem:** Simulation phase is costly

Predicts **win rate** given a board state, avoiding the need of simulating



# Network architecture

- The two networks are actually just one network, with two different heads at the end
- 9 Residual blocks, making wide use of Convolutional and BatchNorm Layers, and of Skip Connections
- 3.1 million parameters in total





# Training

Extract transitions from past years' games:

1

$$\{(board_i, move_i, win/lose_i)\}_{i=1,\dots,N}$$

with  $N = 4971$ . If we use augmentation to exploit symmetry:  $N = 4971 * 8 \approx 40K$

Train the neural network in order to predict:

2

- $board \rightarrow move$  for the Policy Head
- $board \rightarrow win/lose$  for the Value Head

Actually, we are training two different networks, one for the white player and one for the black one, since the game is **asymmetric**. Both are used during the search, alternately.



# Closing remarks

## Limitations:

- In game performance is really poor, with the agent playing almost randomly
- Limited amount of data with respect to the number of parameters (and unknown quality of it) → collect more data
- Probably more time/computing power needed to be able to see learning converge to a competitive player (AlphaGo Zero trained for days on 64 GPUs...)

**Possible improvements:** once a decent starting player is reached, can continue training using self-play

**Bottom line:** what the agent has learned? Not that much. What I've learned? A lot! 😊



Thank you for the attention!